Assignment - 4

Details:

- Assignment Date : 20 October 2022
- Student Name : VENKATA SATWIK RACHAPUTI
- Student Roll Number : 111519106128
- Maximum Marks : 2 Marks

```python
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from keras.models import Model
from keras.layers import LSTM, Activation, Dense, Dropout, Input, Embedding
from keras.optimizers import RMSprop
from keras.preprocessing.text import Tokenizer
from keras.preprocessing import sequence
from keras.utils import pad_sequences
from keras.utils import to_categorical
from keras.callbacks import EarlyStopping
```

READING DATASET

```python
df = pd.read_csv('spam.csv',delimiter=',',encoding='latin-1')
df.head()
```

	v1	v2	Unnamed: 2	Unnamed: 3	Unnamed: 4
0	ham	Go until jurong point, crazy.. Available only ...	NaN	NaN	NaN
1	ham	Ok lar... Joking wif u oni...	NaN	NaN	NaN
2	spam	Free entry in 2 a wkly comp to win FA Cup fina...	NaN	NaN	NaN
3		U dun say so early hor... U c already	NaN	NaN	NaN

```python
df.drop(['Unnamed: 2', 'Unnamed: 3', 'Unnamed: 4'],axis=1,inplace=True)
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5572 entries, 0 to 5571
Data columns (total 2 columns):
 #   Column  Non-Null Count  Dtype
---  ------  --------------  -----
 0   v1      5572 non-null   object
 1   v2      5572 non-null   object
```

● ✕

```
       memory usage: 87.2+ KB
```

```
df.groupby(['v1']).size()
```

```
       v1
       ham      4825
       spam      747
       dtype: int64
```

```
df.groupby(['v2']).size()
```

```
       v2
        &lt;#&gt;   in mca. But not conform.
       1
        &lt;#&gt;   mins but i had to stop somewhere first.
       1
        &lt;DECIMAL&gt; m but its not a common car here so its better to buy from
       china or asia. Or if i find it less expensive. I.ll holla     1
        and   picking them up from various points
       1
        came to look at the flat, seems ok, in his 50s? * Is away alot wiv work.
       Got woman coming at 6.30 too.                                  1

       ..
       ÌÏ still got lessons?   ÌÏ in sch?
       1
       ÌÏ takin linear algebra today?
       1
       ÌÏ thk of wat to eat tonight.
       1
       ÌÏ v ma fan...
       1
       ÌÏ wait 4 me in sch i finish ard 5..
       1
       Length: 5169, dtype: int64
```

```
X = df.v2
Y = df.v1
le = LabelEncoder()
Y = le.fit_transform(Y)
Y = Y.reshape(-1,1)
```

```
X_train,X_test,Y_train,Y_test = train_test_split(X,Y,test_size=0.15)
```

```
max_words = 1000
max_len = 150
tok = Tokenizer(num_words=max_words)
tok.fit_on_texts(X_train)
sequences = tok.texts_to_sequences(X_train)
sequences_matrix = pad_sequences(sequences,maxlen=max_len)
```

```
sequence_matrix = pad_sequences(sequences,maxlen=max_len)
```

CREATE MODEL AND ADD LAYERS

```python
inputs = Input(name='inputs',shape=[max_len])
layer = Embedding(max_words,50,input_length=max_len)(inputs)
layer = LSTM(64)(layer)
layer = Dense(256,name='FC1')(layer)
layer = Activation('relu')(layer)
layer = Dropout(0.5)(layer)
layer = Dense(1,name='out_layer')(layer)
layer = Activation('sigmoid')(layer)
model = Model(inputs=inputs,outputs=layer)
```

COMPILE AND FIT THE MODEL

```python
model.summary()

model.compile(loss='binary_crossentropy',optimizer=RMSprop(),metrics=['accuracy'
model.fit(sequences_matrix,Y_train,batch_size=128,epochs=10,
          validation_split=0.2)
```

```
Model: "model"

_____
 Layer (type)                Output Shape              Param #
=================================================================
 inputs (InputLayer)         [(None, 150)]             0

 embedding (Embedding)       (None, 150, 50)           50000

 lstm (LSTM)                 (None, 64)                29440

 FC1 (Dense)                 (None, 256)               16640

 activation (Activation)     (None, 256)               0

 dropout (Dropout)           (None, 256)               0

 out_layer (Dense)           (None, 1)                 257

 activation_1 (Activation)   (None, 1)                 0

=================================================================
Total params: 96,337
Trainable params: 96,337
Non-trainable params: 0
_____
Epoch 1/10
30/30 [==============================] - 11s 286ms/step - loss: 0.3204 - ac
Epoch 2/10
30/30 [==============================] - 8s 260ms/step - loss: 0.0889 - acc
Epoch 3/10
30/30 [==============================] - 8s 263ms/step - loss: 0.0482 - acc
Epoch 4/10
30/30 [                              ]   - 8s 261ms/step - loss: 0.0361 - acc
```

```
30/30 [==============================] - 8s 261ms/step - loss: 0.0361 - acc
Epoch 5/10
30/30 [==============================] - 8s 258ms/step - loss: 0.0312 - acc
Epoch 6/10
30/30 [==============================] - 10s 328ms/step - loss: 0.0223 - ac
Epoch 7/10
30/30 [==============================] - 9s 284ms/step - loss: 0.0179 - acc
Epoch 8/10
30/30 [==============================] - 10s 344ms/step - loss: 0.0123 - ac
Epoch 9/10
30/30 [==============================] - 8s 271ms/step - loss: 0.0087 - acc
Epoch 10/10
30/30 [==============================] - 8s 262ms/step - loss: 0.0082 - acc
<keras.callbacks.History at 0x7f984f25b450>
```

SAVING THE MODEL

```
model.save('sms_classifier.h5')
```

TEST THE MODEL

```
test_sequences = tok.texts_to_sequences(X_test)
test_sequences_matrix = pad_sequences(test_sequences,maxlen=max_len)
```

```
accr = model.evaluate(test_sequences_matrix,Y_test)
```

```
27/27 [==============================] - 1s 24ms/step - loss: 0.0512 - accu
```

```
print('Test set\n  Loss: {:0.3f}\n  Accuracy: {:0.3f}'.format(accr[0],accr[1]))
```

```
Test set
  Loss: 0.051
  Accuracy: 0.986
```

Colab paid products - Cancel contracts here